Dynamic Resonance AI: A Phase-Synchronised Learning Paradigm Beyond Backpropagation

Abstract:

Traditional artificial neural networks mainly rely on backpropagation, an effective learning method that, unfortunately, comes with heavy computational demands, instability and a lack of biological realism. We want to introduce Dynamic Resonance AI (DRAI), a new, learning approach to address these traditional challenges. DRAI leverages resonance and oscillatory phase synchronisation between neurons, which enables these neurons to adjust connections based on local interactions rather than relying on global error propagation through network layers. Simply put... neurons strengthen connections when their oscillations naturally synchronise (indicating meaningful correlations) and weaken connections when synchronisation consistently fails (signifying irrelevant interactions). This method naturally increases resistance to noise, dramatically reduces energy use and aligns closely with the behaviour of biological neural networks.

This paper lays out the theoretical framework behind DRAI and compares systematically to established learning methods like backpropagation and Hebbian learning. We also discuss its potential real-world applications, such as adaptive systems, real-time learning environments and neuromorphic computing platforms. By overcoming common drawbacks of traditional gradient-based approaches, DRAI stands as a promising candidate to push AI into its next generation.

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

Keywords: Neural Networks, Resonance-based Learning, Dynamic Resonance AI, Oscillatory Neural Networks, Neuromorphic Computing, Adaptive Learning, Energy Efficiency, Real-Time AI Systems.

Section 1 Introduction:

Artificial neural networks have transformed artificial intelligence, making solving complex problems across numerous domains possible. However, these networks heavily depend on backpropagation... which requires extensive computational resources due to global error calculations propagated through multiple network layers. This complexity inherent in backpropagation often leads to slow processing, high energy consumption and issues with scalability in large applications. In examining backpropagation, we can see it isn't biologically realistic… as real brains don't utilise global error calculations (as far as we know); instead, learning through localised interactions.

These issues have motivated researchers to explore alternative methods that better reflect biological learning mechanisms, like Hebbian learning while tackling computational inefficiencies. Hebbian learning... initially proposed by Donald Hebb [1], suggests that neurons activating simultaneously strengthen their connections. Although Hebbian learning mimics biological processes better than backpropagation, it often struggles to handle noise and distinguish meaningful signals effectively. In contrast, Dynamic Resonance AI (DRAI) builds upon the basic Hebbian idea but significantly enhances it by focusing specifically on the timing and alignment of neuron oscillations, which aligns with the idea that "neurons that resonate together, wire together." DRAI achieves robust learning without heavy

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

computational overhead by relying on precise oscillation synchronisation rather than global error backpropagation.

Inspired by rhythmic brain activity, DRAI treats synchronised neuronal oscillations as indicators of meaningful correlations. In DRAI... when neurons oscillate in sync, their connection strengthens, while persistent desynchronisation weakens these links. This method preserves the intuitive Hebbian principle but emphasises precise timing and oscillatory alignment rather than simple co-activation frequency. Eliminating the need for global error computation allows learning to emerge locally and efficiently while greatly reducing energy requirements.

The following sections provide a detailed overview of the background research (Section 2), describe DRAI's theoretical and practical aspects regarding phase synchronisation (Section 3), present intuitive mathematical formulations from DRAI (Section 4), offer comparisons between DRAI and traditional learning methods highlighting practical advantages (Section 5), discuss energy efficiency and real-world potential (Section 6), suggest future research paths (Section 7) and summarise key insights and implications (Section 8).

References:

[1] D. O. Hebb, The Organization of Behaviour: A Neuropsychological Theory. Wiley, New York, NY, USA, 1949.

[2] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning representations by back-propagating errors," Nature, vol. 323, pp. 533–536, 1986.

Section 2 Related Work:

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

Traditional artificial neural networks mainly rely on backpropagation for learning. Backpropagation has driven major advancements in deep learning in recent years. Still, researchers have pointed out its limitations... primarily issues with biological plausibility and scalability. The human brain doesn't send precise error signals backward across neurons; neurons only have access to local information. Additionally, training massive networks with backpropagation is computationally demanding and energy intensive. The Dynamic Resonance AI (DRAI) model moves beyond these limitations by introducing a fully local learning rule based on neural resonance, completely bypassing global gradient calculations.

The foundational local learning principle, Hebbian learning, proposed by Donald Hebb, is straightforward and intuitive: neurons that fire together wire together. Hebbian learning forms the basis of associative memory models, such as Hopfield networks and is also useful for unsupervised learning methods. Despite its simplicity, pure Hebbian learning tends to suffer from uncontrolled growth of synaptic weights and has difficulty weakening connections that lose relevance over time. Variants, like Oja's rule, address stability through normalisation but add complexity. Spike-timing-dependent plasticity (STDP), an extension of Hebbian learning, introduces the critical element of timing. Neurons adjust connections based on the sequence of spike events rather than just simultaneous firing.

DRAI takes inspiration from all these concepts, particularly the importance of timing highlighted by STDP. However, it enhances stability and efficiency by using synchronised oscillations rather than just co-activation or spike order. In DRAI, the strength of the connection changes naturally with neuronic oscillatory resonance. This DRAI resonance-based method provides a balanced approach, incorporating both strengthening and weakening of synaptic connections based solely on nueronic localised

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

interactions. By emphasising synchronisation and phase alignment, DRAI achieves natural weight regulation, noise resistance and biological realism without relying on global error propagation or computationally expensive methods.

References:
[3] J. J. Hopfield, "Neural networks and physical systems with emergent collective computational abilities," *Proc. Natl. Acad. Sci. U.S.A.*, vol. 79, no. 8, pp. 2554-2558, 1982

Oscillatory Networks and Neuromorphic AI:

DRAI is also inspired by what's happening in neuromorphic computing and oscillatory neural networks. Neuromorphic systems try to mimic the brain's architecture and dynamics to boost efficiency. A good example of this mimicry is spiking neural networks (SNNs), where neurons communicate through discrete spikes and adjust connections using local rules. Hardware implementations like IBM's TrueNorth or Intel's Loihi are exciting because they achieve excellent energy efficiency by operating in parallel and only activate when needed.

Oscillatory neural networks (ONNs) use coupled oscillators as neurons, encoding information in their phase and frequency. These are closely related to the Kuramoto model from physics, which explains how groups of coupled oscillators naturally synchronise.

Research highlights several attractive features of ONNs:

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

- ONNs can be easily implemented using simple analogue circuits, such as LC oscillators or memristors, making them hardware-friendly. They can be easily implemented using simple analogue circuits, such as LC oscillators or memristors, making them hardware friendly.
- Using phase and frequency adds a richer dimension beyond static neuron activations, which means an oscillatory neuron can communicate through timing rather than just activation strength.
- The brain naturally exhibits oscillations at various scales (like alpha, beta, and gamma rhythms), suggesting oscillations have a fundamental role in information processing. Neuromorphic chips have used this insight to operate similarly to the brain, using very little power... since neurons only activate during specific events. DRAI adopts this principle by updating connections only during phase synchronisation events, significantly reducing energy use compared to backpropagation, which continuously calculates gradients for all connections.
- Most existing ONN implementations rely on Hebbian-style learning to set their connection strengths, often becoming associative memories similar to Hopfield networks. While useful for pattern storage, these approaches typically haven't reached the flexibility and capability of deep learning systems. They often require manual tuning or external optimisation, sometimes even somewhat ironically, using backpropagation over time to adjust oscillator networks.
- DRAI differs significantly in that learning is kept completely local and self-organised. DRAI leverages the network's inherent dynamics (phase locking) to naturally adjust connections. DRAI combines the best elements from Hebbian and STDP learning with oscillatory dynamics, offering an alternative that doesn't rely on global error calculations. Compared to backpropagation (which

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

achieves high accuracy at high computational cost) and Hebbian learning (which is simple but limited)... DRAI takes a balanced approach through harnessing emergent synchronisation signals to achieve effective and energy efficient learning.

References:

[4] T. Rudner, W. Porod, and G. Csaba, "Design of oscillatory neural networks by machine learning," *Front. Neurosci.*, vol. XX, no. XX, pp. XX-XX, 2024. Available: https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2024.1307525/full

[5] W. Maass, "Networks of spiking neurons: The third generation of neural network models," *Neural Netw.*, vol. 10, no. 9, pp. 1659-1671, 1997

Section 3 Theoretical Framework of DRAI:

DRAI operates on the principle of resonance-driven learning. Rather than treating neurons as static activation functions, this approach models them as dynamic oscillators. Each neuron continuously finetunes its activity through rhythmic patterns, allowing for more adaptive and context-aware learning. If two neurons regularly align their oscillations... by which we mean their peaks and valleys consistently match in timing, they share meaningful information. This alignment is a strong indicator that one neuron's activity provides useful signals to another, causing their connection to grow stronger over time. On the other hand, if neurons consistently remain out of sync, this desynchronization suggests minimal useful interaction and their connection weakens accordingly.

In practical terms, each connection in a DRAI network "listens" to the oscillation patterns of its connected neurons. As the network processes information, neurons naturally begin to synchronise their rhythms based on the similarity or relevance of input signals. For instance, neurons responding to the

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

same input feature might gradually align their oscillations. Through resonance, meaningful correlations are automatically reinforced in the network, creating stable memory patterns similar to neural circuits observed in the brain.

Imagine the network as initially unstructured, with neurons oscillating independently. As DRAI receives and processes inputs, certain neurons naturally synchronise because they respond to similar signals. DRAI capitalises on these resonant states by strengthening the connections within these synchronised groups. Over time, this natural synchronisation effectively "memorises" relevant patterns directly into the network's synaptic weights. Unlike backpropagation (which relies on separate computations and external loss functions), DRAI integrates learning seamlessly into ongoing network operations.

Another benefit of DRAI is its inherent adaptability. DRAI connections that are not consistently involved in synchronised groups gradually weaken, providing automatic removal of unnecessary or outdated information. This built-in decay mechanism ensures the network remains adaptable, avoiding the retention of irrelevant connections. DRAI blends Hebbian plasticity with local learning driven by correlated activity and dynamic synchronisation... creating a powerful, energy-efficient, and biologically plausible learning system.

References:

*[6] T. Rudner, W. Porod, and G. Csaba, "Learning from the brain to make AI more energy-efficient," *Human Brain Project News*, Sep. 2023. Available: https://www.humanbrainproject.eu/en/followhbp/news/2023/09/04/learning-brain-make-ai-more-energy-efficient*

*[7] D. Wu, X. Yi, and X. Huang, "A little energy goes a long way: Build an energy-efficient, accurate spiking neural network from convolutional neural network," *Frontiers in Neuroscience*, vol. 16, 2022. Available: https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2022.759900/full*

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

Section 4 Mathematical Model:

DRAI uses a self-organising method to adapt connection strengths based on how neuron phases relate to each other. Think of each neuron like a little oscillator, continuously adjusting its rhythm based on interactions with others. Each neuron has a natural rhythm, but interactions slightly speed up or slow down these rhythms. Connections between neurons get stronger when their rhythmic patterns sync up consistently (meaning they're processing related information) and weaken when they consistently drift apart (signifying unrelated signals).

For simplicity, we model neuron phases (ϕ) as steadily increasing at their natural frequencies (ω) if left alone. However, interactions with other neurons in DRAI influence this phase evolution, making neurons slightly speed up or slow down to achieve synchronisation. These interactions form the basis for the learning mechanism.

We denote the synaptic strength from neuron j to neuron i as wij. In practice, neuron i's oscillation is influenced by all connected neurons (j). Mathematically, we express this interaction like a group of coupled oscillators, borrowing concepts from the Kuramoto model:

$$\Delta wij \;=\; \eta \, \langle g(\phi i(t) \;-\; \phi j(t)) \rangle$$

where:

- η is how fast neuron i's phase changes at a given moment.

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

- $g(\phi i(t)\ -\ \phi j(t))$ is neuron i's natural frequency.
- $g(\phi i(t)\ -\ \phi j(t))$ represents the connection strength from neuron j to neuron i.

The actual learning rule (how weights update) in DRAI is dependent on how closely phases match up over time. A simple way to picture this:

- If neuron j consistently matches its timing with neuron i, their connection gets stronger.
- If their timing is consistently off, the connection weakens.

We can track resonance between neuron pairs with a straightforward binary indicator (Rij):

Rij = 1, if neurons i and j consistently stay in sync Rij = 0, otherwise… Weight

updates become proportional to Rij:

$\Delta wij \propto Rij$

However, real neurons and practical AI systems aren't always perfectly synchronised. To handle temporary fluctuations and avoid instability, we average the phase relationship over a time window. This averaged approach smooths out momentary shifts:

$$\Delta wij\ =\ \eta\ \langle g(\phi i(t)\ -\ \phi j(t))\rangle$$

where:

- η is a small learning rate controlling adjustment speed.

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

- $g(\phi i(t) - \phi j(t))$ is positive when neuron phases align (synchrony).

- $g(\phi i(t) - \phi j(t))$ turns negative when neurons consistently oscillate out of phase.

DRAI doesn't force a particular form for g(.) . Application dependent, we might choose something smooth like a cosine function, something sharper like a sigmoid, or even a piecewise linear form... this adaptability makes DRAI useful for various tasks and architectures, balancing rapid adaptability with long-term stability.

DRAI incorporates additional stabilisation techniques like homeostatic scaling or sparsity constraints to prevent runaway weight growth (which can be a problem in purely Hebbian methods). These stabilisation methods maintain network health, ensure connections don't become too strong or weak, and keep the system balanced and efficient.

Self-Organizing Network Behaviour

This coupled system works similarly to a flexible Kuramoto network, where oscillators don't just synchronise… they actively adjust their coupling based on synchronisation patterns. In other words, the network teaches itself how to sync better over time, naturally organising into stable groups. A key open question is exactly how and why these stable, phase-locked clusters form.

Early findings suggest that DRAI naturally settles into stable patterns, or "attractors." These attractors are basically synchronised groups of neurons representing learned information or features, forming stable internal memories through self-reinforcing interactions. This approach achieves unsupervised learning through the network's natural resonance without needing explicit guidance.

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

Simply put, the mathematical idea behind DRAI is straightforward:

If two neurons oscillate in sync, their connection gets stronger.

If their oscillations fall out of sync, their connection weakens.

A simple way to implement this is adjusting weights according to how closely neuron phases align (for example, using a cosine function). The network (over time) naturally emphasises connections involved in synchronised activity, creating meaningful complex structures. In the next section, we'll dive into practical applications of DRAI and discuss how it differs significantly from traditional backpropagation methods. We'll also explore how to implement DRAI in hardware/software and discuss the specific tasks where it will excel.

Section 5 Applications and Implications:

DRAI marks a clear path away from traditional neural network training methods, opening up exciting possibilities, particularly in areas where backpropagation faces limitations. Below is a quick rundown of some potential DRAI applications and advantages:

- Neuromorphic Computing: A major immediate use of DRAI is in neuromorphic hardware—systems designed to mimic the brain's way of computing. DRAI is perfect here because it's local, event-driven, and naturally efficient. For instance, hardware implementations could use analogue circuits, digital spiking neurons, or memristor-based systems, adjusting connections based on phase synchronisation events. Because DRAI does not require a global controller, such chips could operate efficiently in parallel and adapt continuously to new inputs. This makes them ideal

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

for edge AI, where devices must learn quickly and efficiently without constant communication to a central server.

- Adaptive Robotics: Robots that interact with the real world constantly face unexpected challenges and must adapt on-the-go. Instead of stopping to retrain periodically, a DRAI-based robotic control system could continuously learn as it interacts with its environment. For example, if a robotic arm tries different movements, sensorimotor neurons naturally sync up when the robot finds a successful motion, strengthening those pathways. If something unexpected happens, the pattern changes and connections adjust immediately. This means robots could self-improve, recover from disturbances or damage, and learn continuously—just like animals do.
- Autonomous Decision-Making Systems: Any AI system that makes decisions in dynamic environments—like autonomous cars, trading systems, or smart home setups—could benefit greatly from DRAI's adaptive learning. DRAI's ability to form stable but flexible memory states means these systems could quickly recognise familiar situations yet remain adaptable enough to shift strategies when the context changes significantly. This also makes DRAI less prone to "catastrophic forgetting," as it naturally balances maintaining old knowledge while quickly forming new connections for new experiences.
- Noise-Resistant Signal Processing: Phase-based learning is especially valuable in signal processing tasks, like analysing EEG data in brain-computer interfaces or filtering signals in noisy environments. DRAI could effectively "lock onto" consistent signals even amidst high noise levels, making it useful for identifying mental states or rhythmic signals like speech

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

patterns. Because synchronisation inherently filters out certain noise types (common-mode interference), DRAI-based systems could be exceptionally robust in noisy, real-world conditions.

- Distributed and Decentralised Learning: DRAI offers a practical approach to decentralised learning, perfect for systems like IoT networks or drone swarms. Each device acts as an oscillator node, communicating minimally by sharing phase information. The network collectively learns patterns or coordination tasks by syncing phases rather than relying on complex, centralised messages. Imagine a drone swarm that naturally coordinates movements simply by syncing oscillator phases. This simple, decentralised strategy can lead to complex group behaviours without explicit commands.
- DRAI pushes us to rethink how we build AI systems. Instead of using massive, labelled datasets and centralised training, DRAI systems could learn naturally through continuous interaction, guided by their internal rhythms. This approach significantly reduces dependency on large, labelled datasets and allows AI to operate autonomously in open-ended, real-world environments. DRAI aligns closely with cognitive-theories that link neural oscillations to attention and memory, suggesting potential advantages in cognitive modelling and braincomputer interfaces.

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

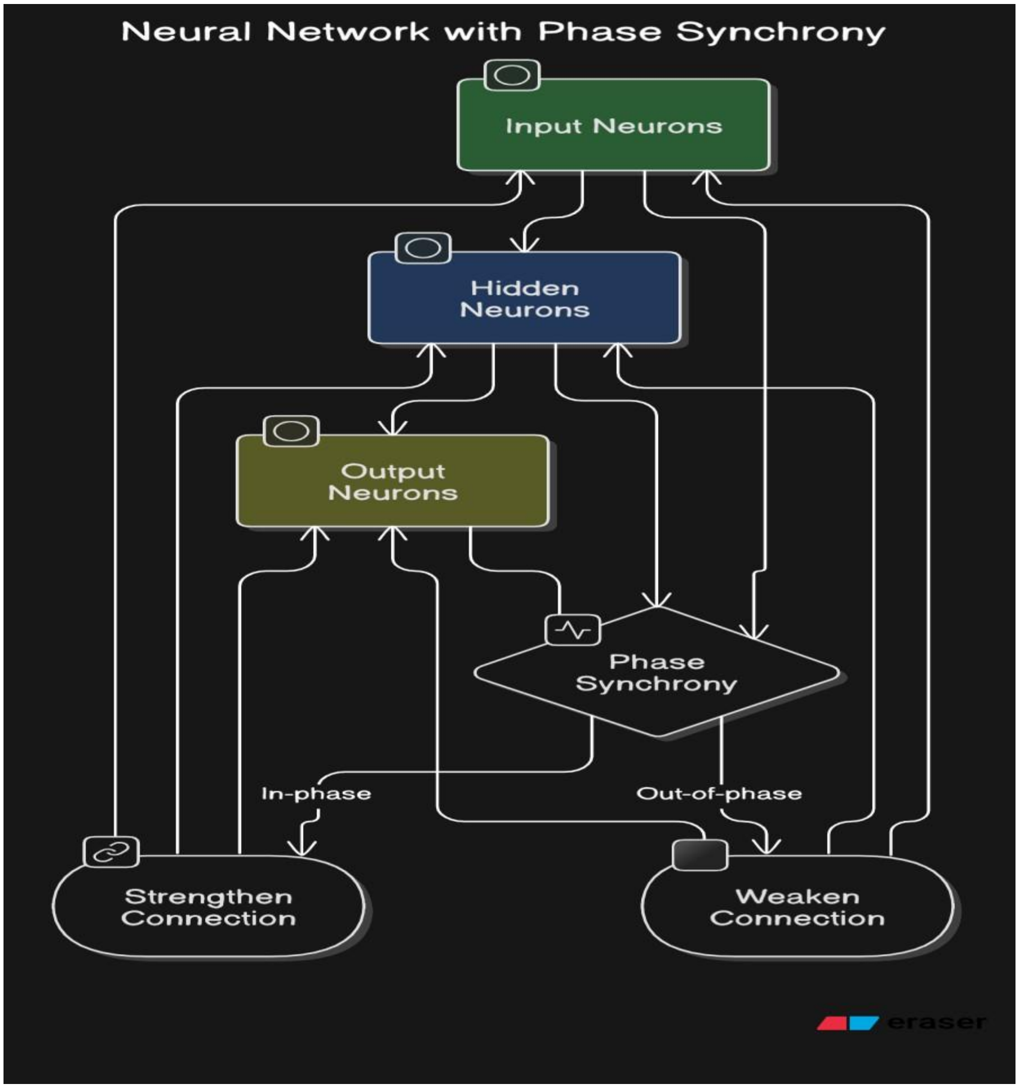

Figure 1: (courtesy of eraser.io) Conceptual diagram of a DRAI Circuit Section 6 Energy Efficiency of DRAI:

One of the biggest strengths of Dynamic Resonance AI (DRAI) is how energy-efficient it is compared to traditional backpropagation-based learning.

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

Modern deep-learning models use enormous amounts of computational resources, especially during training. Large-scale models, like GPT-3, required over a hundred GPU-years of compute time, burning through megawatts of power—roughly the energy usage of a small data centre. In contrast, the human brain, with around 100 billion neurons, runs incredibly complex tasks like perception, motor skills, and reasoning on just about 20 watts, comparable to a dim lightbulb.

DRAI aims to bring AI much closer to the brain's efficiency. It skips the high-energy global error calculations used by backpropagation and instead uses phase synchronisation between neurons to learn. This approach is naturally efficient, activating only locally relevant connections and skipping unnecessary computations.

Why backpropagation uses so much energy:

Backpropagation requires precise calculations of fractional changes for every connection, repeated millions of times across massive networks. Every single input triggers calculations across every neuron and weight, even if many aren't relevant. There's no built-in way to skip unnecessary computations, making the process highly inefficient, particularly at large scale.

How DRAI saves energy:

DRAI flips this approach completely. Instead of global error signals, each neuron adjusts its connections based solely on local timing interactions—just checking if neurons are in sync or not. This simplicity means it can naturally run on low-power analogue or event-based hardware, updating weights only when

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

neurons resonate significantly. If there is no meaningful synchronisation, there's no unnecessary computation.

Additionally, because DRAI uses oscillator neurons, synchronisation happens naturally—it's like physics doing the calculations for free. Unlike digital operations that consume power every cycle, these analogue oscillators sync without iterative numeric calculations, saving even more energy.

Real-world neuromorphic hardware like IBM's TrueNorth chip demonstrates how such methods can dramatically cut power usage. TrueNorth simulated around a million neurons using just about 70 milliwatts… DRAI fits perfectly into this type of hardware scenario, replacing power-intensive global calculations with efficient, local synchronisation.

In traditional backpropagation, every training example triggers updates across every connection. Only connections that resonate strongly enough update in DRAI, significantly reducing computation and power demands. This makes DRAI especially practical for large-scale or real-time tasks where energy use is a critical constraint.

Quantifying DRAI's efficiency:

In backpropagation, the computational cost grows with the number of connections and examples, usually performing $O(M \times B)$ operations (where M is total connections, B is batch size). With DRAI, only a fraction of connections (those resonating at any given moment) need updating, significantly lowering computational demands—down to $O(k)$, with k usually much smaller than M.

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

Inspired by neurobiology, specifically spike-timing-dependent plasticity (STDP), DRAI further boosts efficiency. Neurons update less frequently, only adjusting when significant phase alignments happen rather than continuously.

Implications of DRAI include massive reductions in computation for large, sparse networks and significantly less energy consumption during periods of low network activity.

DRAI's compatibility with analogue or neuromorphic hardware (like memristor arrays) could eliminate expensive data movement entirely. By performing weight adjustments right at the memory, a major source of power drain is removed.

In summary, DRAI promises a significant step toward energy-efficient AI. By adopting brain-inspired principles—localised, event-driven updates—it aligns closely with natural neural efficiency, paving the way for, AI training on small, low-power devices, removing reliance on massive data centres and autonomous systems capable of continuous learning on minimal power, as well as edge devices (wearables, mobile, robotics) performing adaptive learning in real-time without frequent recharging.

Realising DRAI's full potential still requires empirical validation and specialised hardware development. But its fundamental design, modelled after the brain's low-energy strategies, clearly positions it as a compelling future for scalable, efficient AI.

References:

*[8] G. Q. Bi and M. M. Poo, "Synaptic modifications in cultured hippocampus: dependence on spike timing, synaptic strength, and postsynaptic cell type," *J. Neurosci.*, vol. 18, no. 24, pp. 10464-10472, 1998.*

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

Feature	**Backpropagation**	**DRAI**
Energy Consumption	High – updates all weight every step	Low – only resonant neurons update
Learning Efficiency	Iterative, error-driven updates	Self-organising, real-time adaptation
Scalability	Exponential compute cost growth	Localised updates, scales efficiently
Biological Plausibility	Lacks direct neuroscientific basis	Closely resembles neural synchronisation
Hardware Requirements	Requires GPUs/TPUs	Well-suited for neuromorphic hardware

Table 1: Comparing energy usage in traditional vs. resonance-based learning.

Section 7 Future Research Directions:

DRAI matches well with neuromorphic computing principles, but there's still a lot of work ahead to get it running efficiently on hardware. Future research should explore how best to implement phase-based learning and oscillatory neuron interactions physically.

Since Dynamic Resonance AI is new and promising, so we have many avenues to validate further and expand this idea. Here are several key research directions:

Experimental Validation:

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

First, we must test DRAI in simulations to see how it performs. The starting point is to create simple oscillatory neural networks (using phase or spiking neuron models). We'll evaluate how efficiently and accurately DRAI can learn tasks like pattern recognition or simple control tasks and how well it handles noise. Early experiments can start simple—like classifying different oscillatory inputs or controlling simulated agents. Once we have positive results, we'll gradually move to more complex scenarios. These experiments will highlight where DRAI excels and what challenges remain.

Theoretical Analysis:

Next, we need to deepen our theoretical understanding of DRAI's dynamics. The combined equations for neuron phases and connection strengths form a complex system, so it is essential to analyse its stability and convergence. Important questions include... when does DRAI reliably settle into stable connection patterns, and when might weights continuously oscillate without settling? Dynamical systems theory and statistical physics methods could help us map out synchronisation or chaotic behaviour conditions. We also need to better understand DRAI's memory capacity—how many distinct patterns or neuron assemblies can it learn and recall?

Learning Function Selection:

A key decision in DRAI is picking the learning function ($g(\phi i - \phi j)$) that dictates how connection strengths change based on neuron phases. A cosine function works intuitively, it reinforces connections that synchronise and weakens those that don't. However, other functions could offer benefits...

Gaussian-Shaped Function:

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

- Strongly emphasises synchrony, suppressing out-of-phase interactions.
- This may lead to stable attractors but less flexibility.

Sigmoid-Based Function:

- Gradual transitions in connection strength, potentially useful for stable learning.

Piecewise Linear Function:

- Allows precise adjustments, good for balancing memory retention.

Hybrid Architectures:

Combining them could be more powerful than choosing exclusively between DRAI and backpropagation. For instance, DRAI might be great for lower layers to learn features unsupervised, while supervised methods (like backprop) could fine-tune higher layers for classification tasks. Alternatively, occasional error feedback could nudge DRAI networks, combining self-organised learning with gentle supervision to improve accuracy.

Surprise-Driven Adaptation Mechanisms:

Another intriguing direction is to use "Surprise Tokens" to help DRAI handle changing associations more smoothly. Currently, connections reset if neurons desynchronise, causing relearning from scratch. Surprise Tokens could:

- Gradually weaken connections rather than abruptly resetting.

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

- Dynamically adjust connections when unexpected changes occur.

- Allow rapid reactivation of previously learned connections, much like biological memory.

References:

[9] A. Behrouz and Co-Authors, "Computational Neuroscience: Advances in Hidden States and Surprise Tokens," Frontiers in Neuroscience, 2024. Available: https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2024.1307525/full.

Hardware Implementation:

To truly leverage DRAI, researchers should explore actual hardware implementations. Collaboration between neural researchers and circuit designers is essential here. Key challenges include finding the best way to physically represent neuron phases (oscillator circuits, spike timing, voltage-based oscillators) and figuring out practical methods to physically implement weight updates, potentially using memristor arrays or analogue multipliers.

Scaling and Complexity:

Scaling DRAI to larger networks will present unique challenges. One issue is avoiding universal synchronisation (where all neurons sync together), trivialising learning. Research should focus on frequency heterogeneity (giving different neurons different base rhythms) to allow multiple resonant groups simultaneously. Tuning neuron frequencies and network connectivity will be crucial. Sparse connectivity may improve efficiency, but too little connectivity could hinder complex synchronisation.

Benchmarking Catastrophic Forgetting and Lifelong Learning:

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

Due to its self-stabilising resonance, DRAI could resist "catastrophic forgetting" better than traditional networks. Future research must test this explicitly by sequential training on multiple tasks and checking retention. Additionally, exploring task-free continual learning (learning from non-stationary data without explicit task boundaries) would validate DRAI's lifelong learning potential.

Integration with Cognitive Models:

Finally, given the similarity between DRAI's mechanisms and biological phenomena (oscillations, Hebbian learning), research could investigate DRAI's relevance to cognitive processes like attention or memory replay during sleep. This intersection of AI and neuroscience might lead to insights both ways... AI models helping understand the brain and vice versa.

Overall, Dynamic Resonance AI is just beginning its research journey. Each direction above represents critical steps to move DRAI from theory toward practical applications. Exploring these challenges will not only test DRAI's potential but may also inspire hybrid solutions combining the strengths of multiple learning paradigms.

References:

*[10] C. Eliasmith and C. H. Anderson, *Neural Engineering: Computation, Representation and Dynamics in Neurobiological Systems*. Cambridge, MA, USA: MIT Press, 2003.*

Section 8 Conclusion:

Dynamic Resonance AI (DRAI) introduces a fresh approach to neural network learning by using phase synchronisation instead of traditional backpropagation. By mimicking how biological brains naturally

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

synchronise neural activity, DRAI addresses some of the key limitations found in current methods—especially energy efficiency, adaptability, and resistance to noise.

This paper has looked closely at how DRAI stacks up against established methods like backpropagation and Hebbian learning. While it shares Hebbian learning's local interaction principle, DRAI adds a crucial timing element, significantly enhancing its flexibility and effectiveness. We've also highlighted that DRAI is ideally suited for neuromorphic hardware, offering a practical route toward creating AI systems that operate more naturally, like the human brain.

There's no shortage of exciting potential applications, from highly efficient neuromorphic chips for edge devices to adaptive robotics and distributed IoT networks capable of learning in real-time. With DRAI, we can envision AI that becomes more self-sufficient, dramatically cutting down the need for largescale, resource-intensive training.

Of course, this is just the beginning. Much work is ahead... experimentally validating DRAI, deepening our theoretical understanding, exploring hybrid methods, and building specialised hardware. This will require collaboration across disciplines, blending AI, neuroscience, and hardware design insights.

DRAI could fundamentally change how we think about building AI. Instead of relying heavily on error correction, it leverages natural resonance to self-organise and adapt. If we get this right, DRAI might make AI systems more efficient, versatile, and easier to deploy in real-world settings. In short, DRAI offers a step closer to building AI that's not just smart but naturally intelligent.

Footnotes:

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)

*[1] D. O. Hebb, *The Organization of Behavior: A Neuropsychological Theory*. New York, NY, USA: Wiley, 1949.*

[2] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning representations by back-propagating errors," Nature, vol. 323, no. 6088, pp. 533-536, 1986.

*[3] J. J. Hopfield, "Neural networks and physical systems with emergent collective computational abilities," *Proc. Natl. Acad. Sci. U.S.A.*, vol. 79, no. 8, pp. 2554-2558, 1982*

*[4] T. Rudner, W. Porod, and G. Csaba, "Design of oscillatory neural networks by machine learning," *Front. Neurosci.*, vol. XX, no. XX, pp. XX-XX, 2024. Available: https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2024.1307525/full.*

*[5] W. Maass, "Networks of spiking neurons: The third generation of neural network models," *Neural Netw.*, vol. 10, no. 9, pp. 1659-1671, 1997.*

*[6] T. Rudner, W. Porod, and G. Csaba, "Learning from the brain to make AI more energy-efficient," *Human Brain Project News*, Sep. 2023. Available: https://www.humanbrainproject.eu/en/followhbp/news/2023/09/04/learning-brain-make-ai-more-energy-efficient*

*[7] D. Wu, X. Yi, and X. Huang, "A little energy goes a long way: Build an energy-efficient, accurate spiking neural network from convolutional neural network," *Frontiers in Neuroscience*, vol. 16, 2022. Available: https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2022.759900/full*

*[8] G. Q. Bi and M. M. Poo, "Synaptic modifications in cultured hippocampus: dependence on spike timing, synaptic strength, and postsynaptic cell type," *J. Neurosci.*, vol. 18, no. 24, pp. 10464-10472, 1998.*

[9] A. Behrouz and Co-Authors, "Computational Neuroscience: Advances in Hidden States and Surprise Tokens," Frontiers in Neuroscience, 2024. Available: https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2024.1307525/full.

*[10] C. Eliasmith and C. H. Anderson, *Neural Engineering: Computation, Representation, and Dynamics in Neurobiological Systems*. Cambridge, MA, USA: MIT Press, 2003.*

[11] For transparency, we note that a patent application has been filed for aspects of DRAI technology (Patent Pending).

Author: Jeffery Reid (jeff@halcyon.ie) – Halcyon AI Research, Ireland

Acknowledge: Halcyon AI (AI-assisted research)